


QAPINN vs parameter-matched classical twin, identical architecture downstream, identical training

(a) Accuracy, mean $\pm$ s.d. over 5 seeds



(b) Paired ratio with 95% CI
(above parity $\Rightarrow$ QAPINN better)

